

# Composable ML

Boris Arnoux

2016

## Two lessons from the Netflix prize

- The winners used models that combined different ML approaches.
- The winning algorithm was never used in production.

# Context

Relevant problems to this presentation:

- Supervised regression or classification
- Simple structures in feature space (not cats)

Good examples: online recommendations, Ad tech, some finance, some physical models.

# Linear methods rock

- Simple, fast, transparent.
- Online learning available
- Scalable, CPU & memory efficient.
- Sparse

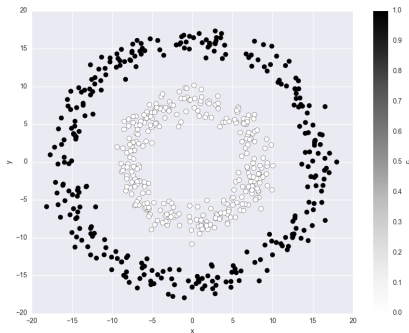
Such practical, operational and computational benefits are very important at scale.

# Non linearity

Not all patterns are linearly separable.


# Non linearity

Not all patterns are linearly separable.



# Non linearity

Not all patterns are linearly separable.



# Kernels

Theoretical justification for the use of kernels in linear models:

$$L(X.\mu, y) \equiv L(XX^T.v, y)$$



Since  $\mu$  is learned against  $X$ ,  $\mu$  can be spanned by  $X$ .

Most famous algorithm: kernel support vector classification.





# What a kernel machine sees (1/2)

L2 Logistic regression, RBF kernel, nonlinear dataset:



# What a kernel machine sees (2/2)

L2 Logistic regression, RBF kernel, arcs dataset:



Here  $\gamma$  (as in  $K(x_1, x_2) = \exp(-\gamma \|x_1 - x_2\|^2)$ ) is chosen to highlight how the new feature space is built. For each sample, the RBF kernel constructs a local indicator variable in the original feature space. Each sample can become such a feature.

# Cogs in the kernel machine

How a typical kernel decides how to classify a new sample  $x_{new}$ :

- $x_{new}$  is compared to the training  $X = (x_1, x_2, \dots, x_n)$  using  $K(., .)$ .
- The sample-to-sample distance  $K(., .)$  makes use of the original feature space.
- The kernel based features values  $k = K(x_i, x_{new})$  are used for computing  $y_{new} = v.k$ .



# Weaknesses of kernels

This leads to two points:

- 1 Often  $K(., x)$  uses all the features in the original space, which as dimensionality grows, eventually scrambles relevant dimensions with the less relevant dimensions.
- 2 Scalability issues (often  $O(N^3)$  complexity)



# Kernels with admixture of irrelevant features

When all features ( $x$ ,  $y$ ) are relevant to the pattern:




# Kernels with admixture of irrelevant features

When two irrelevant features are added (of similar variance):



# Kernels with admixture of irrelevant features

When four irrelevant features are added (of similar variance):



The pure-kernel approach breaks down.

# Kernel machines weaknesses

Mitigating kernel weaknesses (1):

- 1 High number of features: use only a subset of original features, where distances make sense.

But how to bring in new information if we can't use new features?




# Explicit feature maps

Explicit kernel transform step-by-step:

- Pick features from the original feature space which makes sense to include in kernel calculations.
- Choosing "Interesting" samples, to promote as features  
 $f_x(.) = K(., x)$  (lots of ways to be smart here!)
- Augment (rather than replace) the original feature space with these features.


# Explicit feature maps

L2 Logistic regression, "arcs" dataset, decision function for RBF kernel feature transform:




# Explicit feature maps

Decision function when the weights of the original features are erased after training (arcs dataset, RBF feature transform, L2 Logistic regression):



# Kernels with admixture of irrelevant features

Original features augmented with an explicit mapping of kernel features, four irrelevant features added:



In this case, the kernel features break down too, but the model returns to a linear treatment.

# Kernel machines weaknesses

Mitigating kernel weaknesses (2):

- 1 Scalability issues: use kernel approximation.

# Addressing scalability issues

Nystroem sampling:

- Approximates any kernel.
- Based on sampling & interpolation.

# Addressing scalability issues

For linear classification & regression purposes, it is (mostly) equivalent to picking random samples as features as opposed to promoting all the samples.






Figure: A Nystroem dimension.

# Addressing scalability issues



Nystroem sampling in action, L2 Logistic regression, 50 kernel dimensions & 100 samples:





# Addressing scalability issues



Nystroem sampling breaking down, L2 Logistic regression, 10 kernel dimensions & 100 samples:



Important note: vanilla Nystroem sampling is unsupervised and does not attempt to find the best samples to pick.


# Addressing scalability issues

Alternative, Fourier RBF kernel approximation



# Addressing scalability issues

Fourier RBF kernel approximation, zooming out:



# Addressing scalability issues

You can make your own kernel specific approximation, promote samples based on:

- Feature space coverage.
- Where it helps the loss function.

Optionally go through a step of Nystroem or Cholesky for "normalization".

# Kernels wrap up

Kernels allow non linear learning, and explicit mappings allows:

- Properly taking into account new features.
- Gives a lot of engineering latitude in limiting the dimension of the new feature space by sampling.

# Non linearity II

Special cases of non-linearity merit special treatment:

- per-feature internal structure.
- particular features combinations.

# Non linearity II

Special cases of non-linearity merit special treatment:

- per-feature internal structure.
- particular features combinations.

Feature engineering helps with specialized feature transforms:

- Split a feature in binary bins (indicator variables) or linear steps (within-bin barycentric coordinates).
- Bins extend to feature pairs (products of feature bin indicators).

# Non linearity II

How to find these relevant combinations systematically, and optimally?





# Non linearity II

How to find these relevant combinations systematically, and optimally? We need a supervised feature transform.



# Boosting feature transform

Original features augmented with boosting features, L2 Logistic regression:



# Boosting feature transform

Original features augmented with boosting features, L2 Logistic regression:



# Boosting transforms basics

What is a boosting transform:

- Train a gradient boosting model.
- Each leaf of each tree is an indicator variable.
- Augment the initial feature space with the leaf indicators as features.

They are the workhorse of CTR prediction at Facebook (see ADKDD 2014)

## Boosting feature transform analogies:

- one-level tree (decision stump): analogous to feature bin indicator.
- multiple branching levels: analogous to feature tuples.
- minimum samples per leaf: analogous to quantile binning.



# Boosting transforms pros & cons

Boosting transform are great, they are

- Sparse.
- Supervised: mostly no need to worry about irrelevant features.
- A bit rough around the edges...



# Boosting & Mixed classes

L2 Logistic regression, boosting transform alone vs boosting plus RBF features



# Boosting & Mixed classes

L2 Logistic regression, boosting transform alone vs boosting plus RBF features





# Conclusion

An approach based on composable feature transforms, with:

- Linear learning core (with all the benefits)
- Feature transforms create a white box, supervised map of the feature space.
- Feature transforms operate correctly side by side, with other transforms and with linear features.

It contrasts with the "pick the right black box" approach.

Code on github [https://github.com/borithefirst/epfl\\_pres/blob/master/kernel\\_lin.py](https://github.com/borithefirst/epfl_pres/blob/master/kernel_lin.py)